

From Drift-Diffusion Processes to Monte Carlo Sampling

ELG 5218 - Uncertainty Evaluation in Engineering Measurements and Machine Learning

Miodrag Bolic

University of Ottawa

January 28, 2026

Why Drift–Diffusion \rightarrow Langevin \rightarrow Monte Carlo?

Diffusion models taught us a key tool: the *score* (gradient of log density) guides sampling.

Same ingredients, different goal

- **Diffusion models:** learn scores across noise levels (p_t) to generate data.
- **Langevin dynamics:** use the score of a *fixed* target $p(x)$ to generate samples from $p(x)$.

Then: Monte Carlo turns samples into answers

- Once we can sample, we can estimate expectations (LLN) and quantify uncertainty.
- This is the backbone of MCMC: sample from $\pi(\theta)$, then estimate $E[f(\theta)]$.

Next: replace $p(x)$ with Bayesian posterior $\pi(\theta)$ via $U(\theta) = -\log \pi(\theta)$ (MALA).

Step 1: Pure Diffusion — No Preferred Direction

Stochastic Differential Equation (SDE):

$$dX_t = \sigma dW_t \quad (\mu = 0)$$

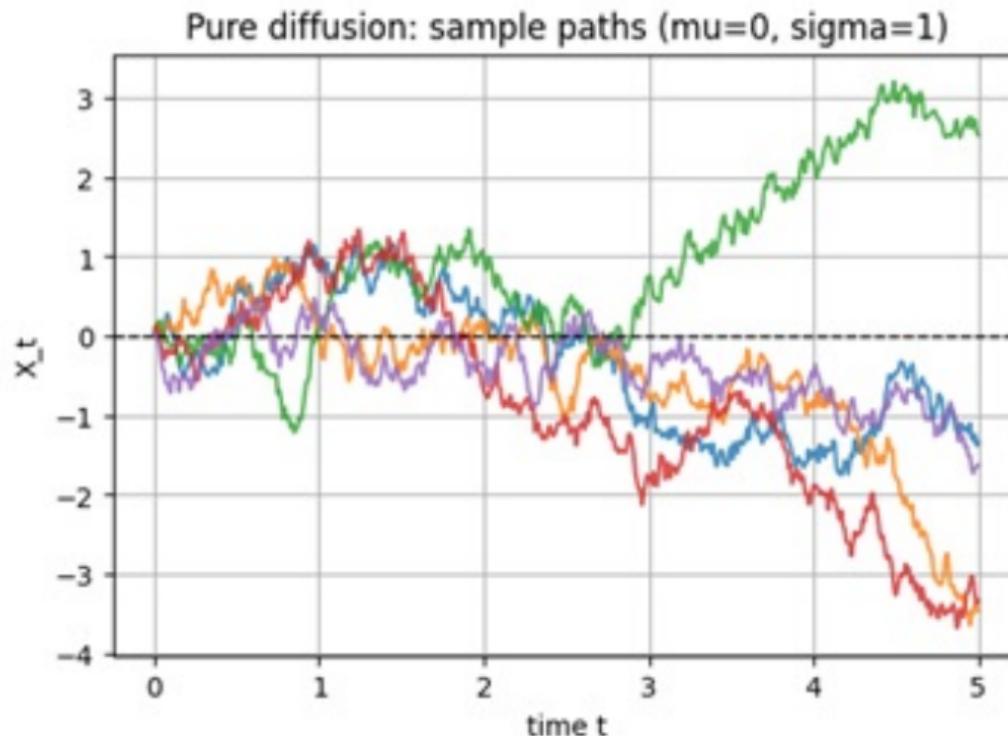
Key Properties:

- **No drift term:** Position wanders randomly with zero mean
- **Symmetric spread:** Distribution stays centered at $X_0 = 0$
- **Variance grows:** $\text{Var}(X_t) = \sigma^2 t$
- **Diffusion:** spreading without a preferred direction
- **Example:** Standard Brownian motion (scaled by σ)

Intuition: Like a particle undergoing random thermal motion with no bias.

Numerical Solution: $X_{k+1} = X_k + \sigma\sqrt{\Delta t} \cdot Z_k, \quad Z_k \sim \mathcal{N}(0, 1)$

Pure Diffusion: Sample Trajectories



Observation: Multiple paths spread out symmetrically around zero.

Step 2: Pure Drift — Deterministic Motion

Stochastic Differential Equation (SDE):

$$dX_t = \mu dt \quad (\sigma = 0)$$

Analytical Solution:

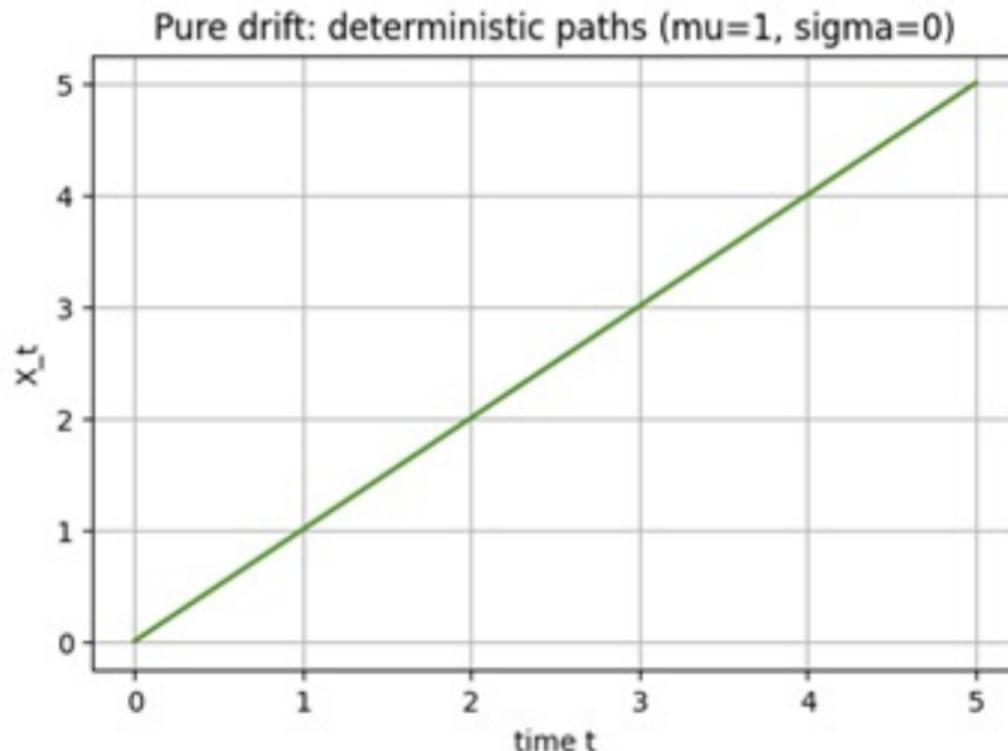
$$X_t = X_0 + \mu t \quad (\text{straight line})$$

Key Properties:

- **Deterministic:** No randomness; same path every time
- **Linear motion:** Moves at constant velocity μ
- **No uncertainty:** All paths identical
- **Example:** $\mu = 1.0$ gives $X_t = t$

Intuition: Like an object rolling downhill with fixed velocity.

Pure Drift: Deterministic Paths



Observation: All paths are identical (deterministic)—no randomness!

Step 3: Drift + Diffusion — Noisy Upward Trend

Stochastic Differential Equation (SDE):

$$dX_t = \mu dt + \sigma dW_t \quad (\mu \neq 0, \sigma \neq 0)$$

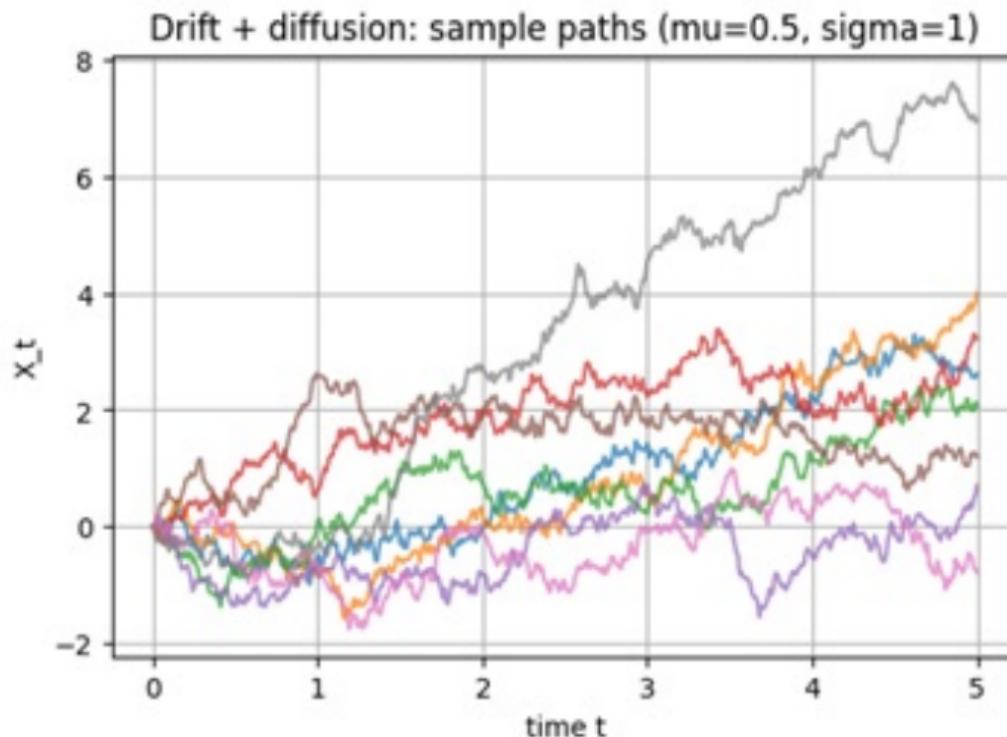
Key Properties:

- **Drift + noise:** Trajectory trends upward on average, with random fluctuations
- **Mean motion:** $\mathbb{E}[X_t] = X_0 + \mu t$ (like pure drift)
- **Growing variance:** $\text{Var}(X_t) = \sigma^2 t$ (like pure diffusion)
- **Example:** $\mu = 0.5$, $\sigma = 1.0$ gives upward trend with noise

Intuition: Like a particle biased downhill but with thermal random kicks.

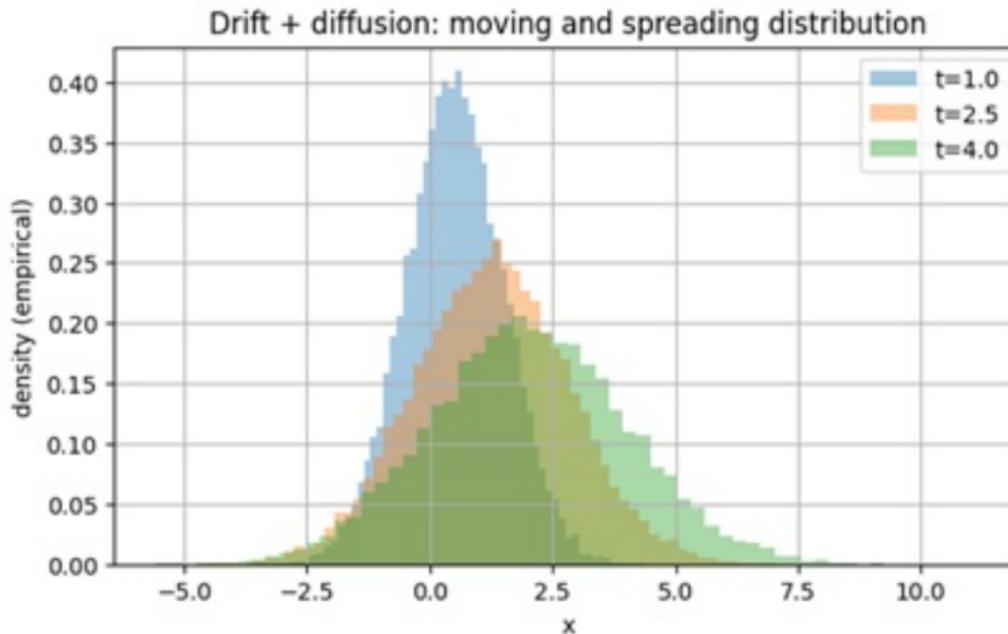
Numerical Solution: $X_{k+1} = X_k + \mu \Delta t + \sigma \sqrt{\Delta t} \cdot Z_k$

Drift + Diffusion: Sample Trajectories



Observation: Paths trend upward on average but with visible random fluctuations.

Drift + Diffusion: Moving and Spreading Distribution



Key Insight: Distribution shifts ($\text{mean} = \mu t$) and spreads ($\text{variance} \propto t$).

Step 4: Langevin Dynamics — Sampling from a Known Distribution

Key Idea: Use drift from gradient of log-target (score) to sample efficiently: $\text{score}(x) = \nabla_x \log p(x)$

Langevin dynamics: a stochastic process that uses this score to move toward high-probability regions, while noise keeps exploration (so we sample rather than optimize).
(drift + Brownian noise)

Langevin SDE:

$$dX_t = \frac{1}{2} \nabla_x \log p(x) dt + dW_t$$

Simple Example — Standard Gaussian:

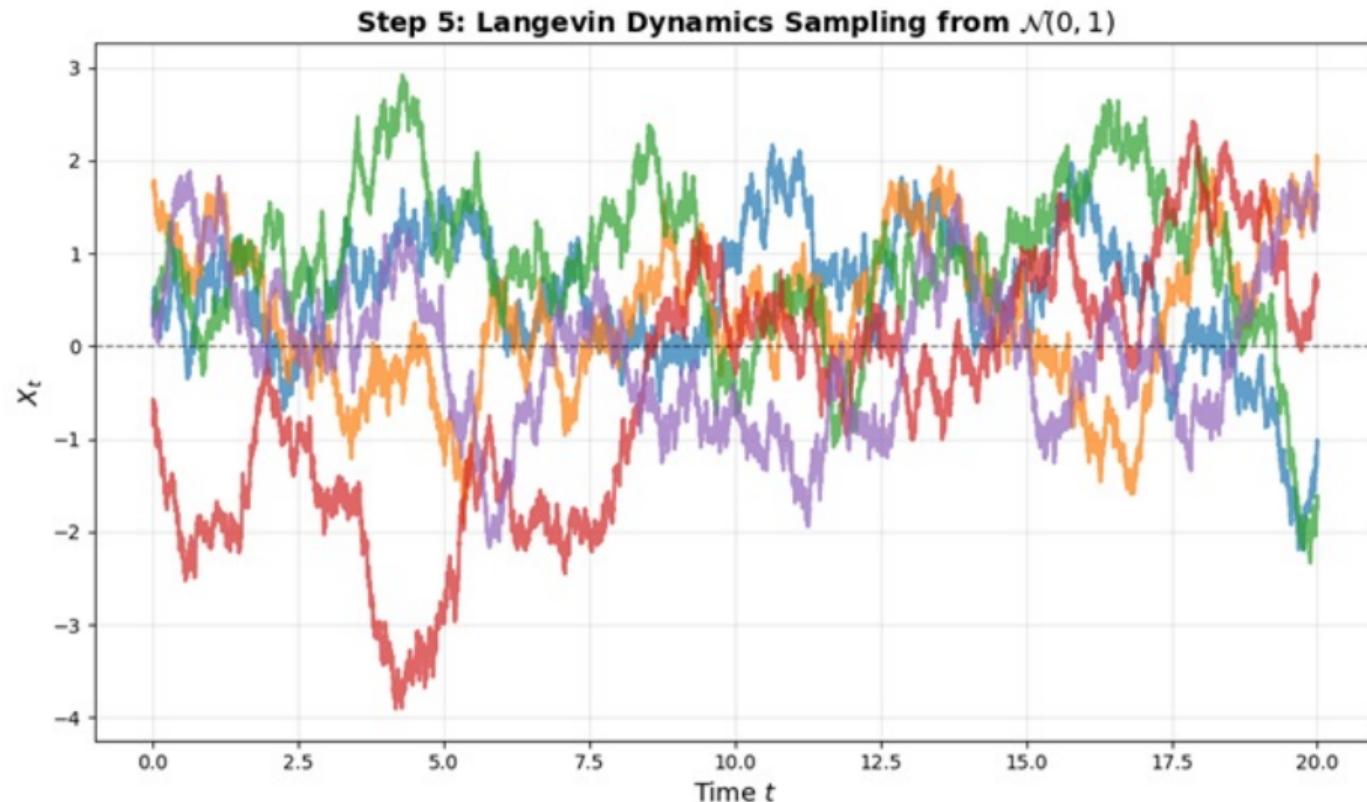
- Target: $p(x) = \mathcal{N}(0, 1)$
- Log-density: $\log p(x) = -\frac{1}{2}x^2 + \text{const}$
- Gradient: $\nabla_x \log p(x) = -x$
- Drift: $b(x) = -\frac{1}{2}x$ (pulls toward zero)

Result: Trajectories concentrate around the mode (here, $x = 0$).

Why It Works:

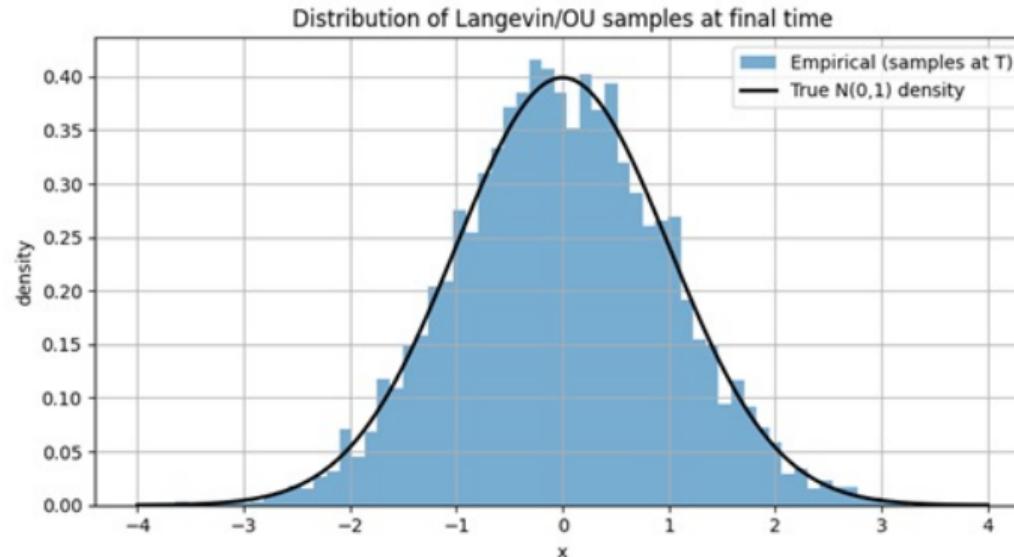
- Drift pulls toward high-probability regions
- Noise prevents convergence to single point
- Stationary distribution = target $p(x)$

Langevin: Trajectories Concentrate Near Mode



Observation: Trajectories spend more time near $x = 0$ (the mode of $\mathcal{N}(0, 1)$).

Langevin: Samples Match Target Distribution



Key Result: Empirical histogram (from Langevin samples) matches $\mathcal{N}(0, 1)$ perfectly!

Step 5: Monte Carlo Estimation — Law of Large Numbers

Problem: Estimate $\mathbb{E}[X_T]$ from many drift-diffusion trajectories.

Setup:

- Generate N independent trajectories of $dX_t = \mu dt + \sigma dW_t$
- Observe value X_T at time T for each trajectory
- Estimate: $\hat{\mathbb{E}}[X_T] = \frac{1}{N} \sum_{n=1}^N X_T^{(n)}$
- True value: $\mathbb{E}[X_T] = X_0 + \mu T$ (by linearity of expectation)

Law of Large Numbers:

$$\hat{\mathbb{E}}[X_T] \xrightarrow{N \rightarrow \infty} \mathbb{E}[X_T] \quad \text{almost surely}$$

Error decreases as: $\text{MSE} \sim \frac{1}{N}$ (independent of dimension!)

Monte Carlo Advantage: Dimension-independent convergence rate, unlike quadrature methods.

From Langevin to MALA: Gradient-Based Bayesian Inference for Complex Posteriors

- ➊ Understand potential energy: $U(\theta) = -\log \pi(\theta)$
- ➋ Learn Langevin dynamics and why it samples from $\pi(\theta)$
- ➌ Understand discretization bias in ULA
- ➍ Master MALA algorithm with Metropolis correction
- ➎ Compare MALA vs Random Walk MH vs HMC
- ➏ Apply to Bayesian logistic regression
- ➐ Verify convergence with diagnostics

Outcome: Gradient-based sampling for any Bayesian model

Step 1: What is $U(\theta)$?

Potential Energy = Energy Landscape of Posterior

Bayes' theorem:

$$\pi(\theta|y) = \frac{p(y|\theta)p(\theta)}{p(y)} \propto p(y|\theta)p(\theta)$$

Define potential energy:

$$U(\theta) = -\log p(y|\theta) - \log p(\theta)$$

Then:

$$\pi(\theta|y) = \frac{1}{Z} \exp(-U(\theta))$$

where $Z = \int \exp(-U(\theta))d\theta$ is the partition function.

Intuition:

- Low U = high posterior density
- Samples concentrate where U is small
- Gradient ∇U points uphill

Step 1 (continued): Logistic Regression

Computing $U(w)$ and $\nabla U(w)$

Model:

$$p(y_i = 1 | x_i, w) = \sigma(x_i^T w) \quad (1)$$

$$p(\beta) = \mathcal{N}(0, \lambda^{-1} I) \quad (2)$$

where $\sigma(z) = \frac{1}{1+e^{-z}}$ is the logistic sigmoid.

Potential energy:

$$U(\beta) = - \sum_{i=1}^n [y_i \log \sigma_i + (1 - y_i) \log(1 - \sigma_i)] + \frac{\lambda}{2} \|w\|_2^2$$

Gradient:

$$\nabla U(w) = X^T(\sigma - y) + \lambda w$$

where $\sigma_i = \sigma(x_i^T w)$.

Cost: $O(nd)$ for gradient computation.

Step 2: Langevin Dynamics

Continuous-Time Sampling via SDE

Langevin SDE (continuous time):

$$dX_t = -\frac{1}{2} \nabla U(X_t) dt + dW_t$$

Key insight: As $t \rightarrow \infty$, the distribution of X_t converges to $\pi(\theta) \propto \exp(-U(\theta))$.

Two components:

- $-\frac{1}{2} \nabla U(X_t) dt$: Gradient descent (push toward mode)
- dW_t : Brownian noise (exploration)
- Balance = samples from correct posterior!

Why this works: the Langevin diffusion is constructed so that its stationary distribution is

$$\pi(\theta) \propto e^{-U(\theta)}.$$

(We use this fact rather than proving it via PDEs in this lecture.)

Step 3: Discretization - ULA

Unadjusted Langevin Algorithm (Biased)

Discrete Euler-Maruyama approximation:

$$\theta_{k+1} = \theta_k - \frac{1}{2} \nabla U(\theta_k) \Delta t + \sqrt{\Delta t} Z_k$$

where $Z_k \sim \mathcal{N}(0, I)$ are i.i.d. standard normals.

Problem: The stationary distribution is **NOT** exactly $\pi(\theta)$.

Discretization bias: The true stationary distribution is

$$\pi_{\Delta t}(\theta) \approx \pi(\theta) + O(\Delta t) \text{ correction terms}$$

Trade-off:

- Small Δt : Less bias, but slow mixing
- Large Δt : Fast mixing, but large bias!

Solution: Add Metropolis acceptance step \Rightarrow MALA (exact!)

Step 4: MALA (Metropolis-Adjusted Langevin)

Exact MCMC via Gradient + Metropolis Correction

Algorithm:

① **Propose:** $\tilde{\theta} = \theta_k - \frac{1}{2}\nabla U(\theta_k)\Delta t + \sqrt{\Delta t} Z_k$

② **Acceptance probability:**

$$\alpha = \min \left(1, \exp \left(-U(\tilde{\theta}) + U(\theta_k) \right) \right)$$

③ **Update:**

$$\theta_{k+1} = \begin{cases} \tilde{\theta} & \text{with probability } \alpha \\ \theta_k & \text{with probability } 1 - \alpha \end{cases}$$

Key properties:

- **Exact:** Stationary distribution is exactly $\pi(\theta)$ (no bias!)
- **Efficient:** Typically 50–70% acceptance rate
- **Fast:** 5–10x faster than random walk MH
- **Cost:** Same $O(nd)$ per iteration

Step 4 (continued): Why This Acceptance Ratio?

General Metropolis-Hastings:

$$\alpha = \min \left(1, \frac{\pi(\tilde{\theta})q(\theta_k|\tilde{\theta})}{\pi(\theta_k)q(\tilde{\theta}|\theta_k)} \right)$$

For Langevin proposal: The proposal is Gaussian

$$q(\tilde{\theta}|\theta_k) = \frac{1}{(2\pi\Delta t)^{d/2}} \exp \left(-\frac{\|\tilde{\theta} - \mu_k\|^2}{2\Delta t} \right)$$

where $\mu_k = \theta_k - \frac{1}{2}\nabla U(\theta_k)\Delta t$.

[Click here for visualization.](#)

Bayesian Logistic Regression with MALA

Data: $X \in \mathbb{R}^{n \times d}$, $y \in \{0, 1\}^n$ (0=functional, 1=failed)

Model:

$$P(y_i = 1 | x_i, w) = \sigma(x_i^\top w), \quad \sigma(z) = \frac{1}{1 + e^{-z}}$$
$$w \sim \mathcal{N}(0, \lambda^{-1} I)$$

Target posterior:

$$\pi(w) \propto p(y | X, w) p(w) \iff U(w) = -\log \pi(w)$$

Why MALA (vs random-walk MH):

- Random-walk MH proposes blindly: $w^* = w + \epsilon$
- MALA proposes using local geometry (a drift term):

$$w^* = w - \frac{\eta}{2} \nabla U(w) + \sqrt{\eta} \xi, \quad \xi \sim \mathcal{N}(0, I)$$

Algorithm Comparison

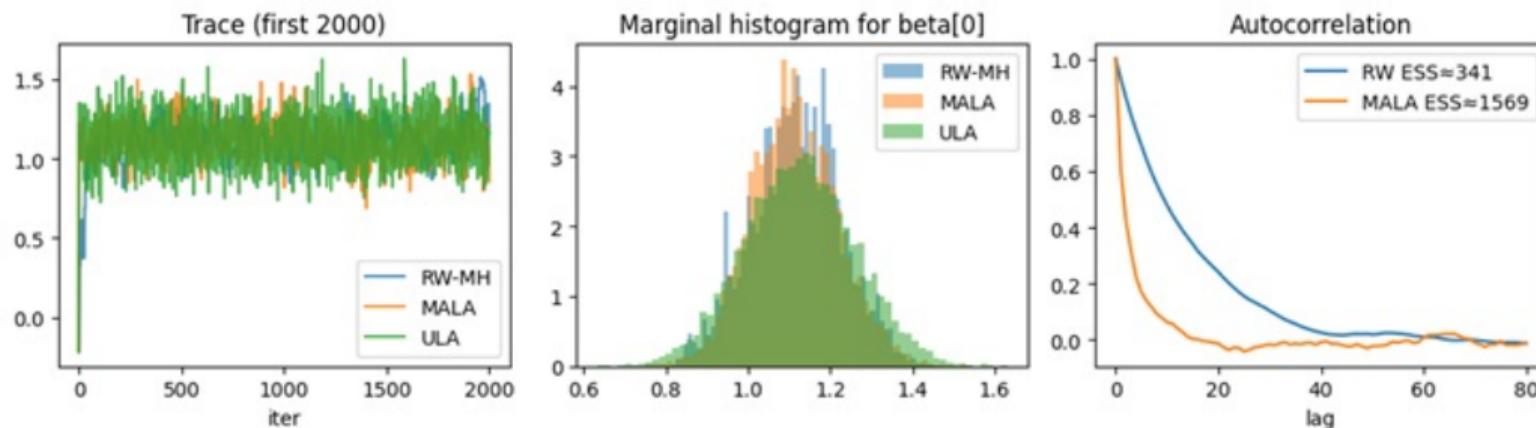


Figure: Comparing MALA and MH for a binary regression problem

Computing Gradients

General template:

$$\nabla U(\theta) = -\nabla \log p(y|\theta) - \nabla \log p(\theta)$$

Common models:

- ① **Linear regression:** $\nabla U = X^T(X\beta - y) + \lambda\beta$
- ② **Logistic regression:** $\nabla U = X^T(\sigma(X\beta) - y) + \lambda\beta$
- ③ **Poisson GLM:** $\nabla U = X^T(\exp(X\beta) - y) + \lambda\beta$
- ④ **Gaussian process:** $\nabla U = (y - \mu)^T K^{-1} \nabla \mu$

Automatic differentiation: Use PyTorch, TensorFlow, or JAX!

Testing: Always verify with finite differences:

$$\frac{\partial U}{\partial \theta_j} \approx \frac{U(\theta + \epsilon e_j) - U(\theta - \epsilon e_j)}{2\epsilon}$$

Checking Your Chains: Diagnostics

1. Trace plots: Plot θ_k vs. iteration k

- Good: Random scatter, no trends
- Bad: Visible trends, stuck values

2. Acceptance rate:

- Good: 50–70% (ideal for MALA)
- Too low: Increase Δt
- Too high: Decrease Δt

3. Autocorrelation function (ACF):

$$\text{ACF}(k) = \frac{\text{Cov}(\theta_t, \theta_{t+k})}{\text{Var}(\theta_t)}$$

Effective sample size: $\text{ESS} = \frac{N}{1 + 2 \sum_{k=1}^{\infty} \rho_k}$

4. Gelman-Rubin statistic (\hat{R}): $\hat{R} = \sqrt{\frac{\text{Var}_{\text{between}}}{\text{Var}_{\text{within}}}}$

- $\hat{R} < 1.01$: Converged
- $\hat{R} > 1.05$: Not converged

Practical Tips for MALA

1. Data preprocessing:

- Standardize: $X = (X - \bar{X})/\sigma_X$
- Avoids numerical issues

2. Tune step size Δt :

- Start with $\Delta t = 0.01$
- Adjust for approximately 60% acceptance
- Test gradient with finite differences first!

3. Burn-in and chains:

- Burn-in: Discard first approximately 25% of iterations
- Run 3–4 independent chains
- Check $\hat{R} < 1.01$ for convergence

4. If things go wrong:

- Low acceptance? Decrease Δt
- Gradient NaN? Standardize data
- Divergent? Use stronger priors
- Slow mixing? Reparameterize model

Summary: From Langevin to MALA

- ① **Potential energy** $U(\theta) = -\log \pi(\theta)$ defines posterior
- ② **Langevin SDE** has stationary distribution $\pi(\theta)$
- ③ **Discretization** introduces bias, Metropolis fixes it
- ④ **MALA:** Propose via Langevin, accept/reject via Metropolis
- ⑤ **Efficiency:** 50–70% acceptance, 10x faster than RW-MH
- ⑥ **Gradient computation:** Essential! Use autodiff
- ⑦ **Diagnostics:** Trace plots, ACF, \hat{R} , ESS
- ⑧ **Applicable:** Any model where you can compute $\nabla U(\theta)$

Gradient-based sampling is the future of Bayesian inference!

Further Reading

- Pavliotis, G.A. (2014). *Stochastic Processes and Applications*. Springer. - Rigorous treatment of SDEs and Fokker-Planck equations
- Girolami, M. & Calderhead, B. (2011). Riemann manifold Langevin and Hamiltonian Monte Carlo methods. *JMLR*. - MALA on manifolds; metric-adaptive proposals
- Roberts, G.O. & Rosenthal, J.S. (2004). General state space Markov chains and MCMC algorithms. *Probability Surveys*. - Theory of convergence and efficiency
- Hoffman, M.D. & Gelman, A. (2014). The No-U-Turn Sampler. *JMLR*. - Modern HMC with adaptive trajectory length
- Betancourt, M. (2017). A Conceptual Introduction to Hamiltonian Monte Carlo. arXiv:1701.02434. - Intuitive explanation of gradient-based sampling
- **Probabilistic Programming:** Stan, PyMC, Turing.jl all use NUTS (refinement of HMC/MALA ideas).

Implementation Resources

For real Bayesian inference, use probabilistic programming:

Tool	Highs
Stan	Automatic NUTS sampler, wide model support
PyMC	Python-native, integrates with numpy/pandas
Turing.jl	Julia (fast), embedded in language
Numpyro	weight, JAX backend, modern autodiff
Blackjax	Composable MCMC kernels, JAX